

# Assignment 3

## Exercise 1

In the lab, we applied random forests to the Boston data using  $mtry = 6$  and using  $ntree = 25$  and  $ntree = 500$ . Create a plot displaying the test error resulting from random forests on this data set for a more comprehensive range of values for  $mtry$  and  $ntree$ . You can model your plot after Figure 8.10. Describe the results obtained.

```
library(MASS)
library(randomForest)

# Construct the train and test matrices
set.seed(123)
train = sample(dim(Boston)[1], 2*dim(Boston)[1]/3, replace = F) # 2/3 split

# Attributes
X.train = Boston[train, -14]
X.test = Boston[-train, -14]

# Response
Y.train = Boston[train, 14]
Y.test = Boston[-train, 14]

# Number of variables used in each split
p = dim(Boston)[2]-1 # All attributes
p.2 = p/2 # Half
p.sq = sqrt(p) # Square root of all possibilities

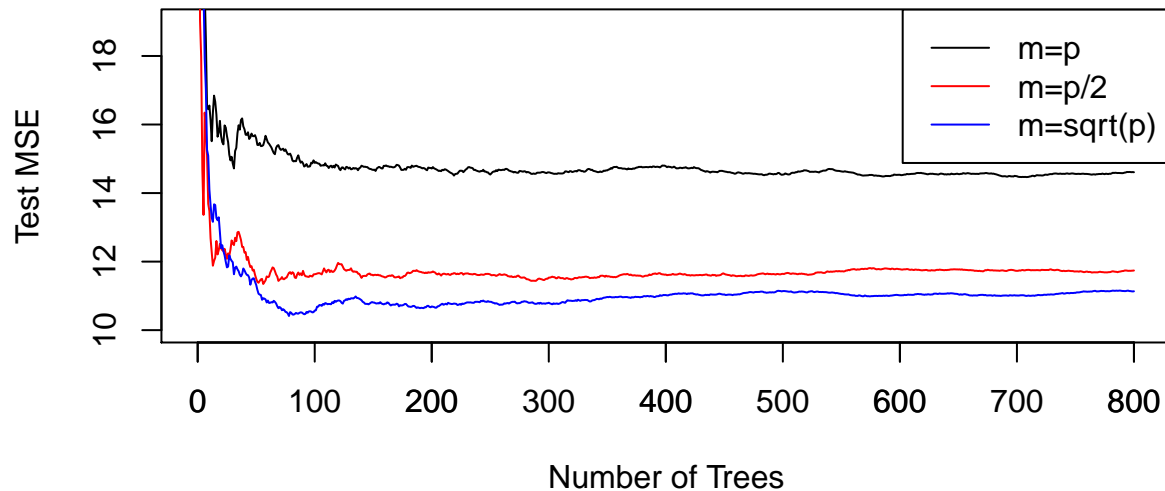
# Models for each split type, number of trees is 800
rf.boston.p = randomForest(X.train, Y.train, xtest=X.test,
                           ytest=Y.test, mtry=p, ntree=800)
rf.boston.p.2 = randomForest(X.train, Y.train, xtest=X.test,
                             ytest=Y.test, mtry=p.2, ntree=800)
rf.boston.p.sq = randomForest(X.train, Y.train, xtest=X.test,
                              ytest=Y.test, mtry=p.sq, ntree=800)

# Plotting the results
plot(1:800, rf.boston.p$test$mse, col=1, type="l", xlab="Number of Trees",
     ylab="Test MSE", ylim=c(floor(min(c(rf.boston.p$test$mse, rf.boston.p.2$test$mse,
rf.boston.p.sq$test$mse))),19), main="Random forest test errors for Boston data")
lines(1:800, rf.boston.p.2$test$mse, col=2, type="l")
lines(1:800, rf.boston.p.sq$test$mse, col=4, type="l")
legend("topright", c("m=p", "m=p/2", "m=sqrt(p)"), col=c(1, "red", "blue"),
      cex=1, lty=1)
axis(1, seq(0,800,100))

paste("The lowest test error:", round(min(c(rf.boston.p$test$mse,
rf.boston.p.2$test$mse,rf.boston.p.sq$test$mse)),2))
```

```
[1] "The lowest test error: 10.41"
```

## Random forest test errors for Boston data



The lowest test set error is achieved with  $\sqrt{p}$  as a criterion for each split, where  $p$  is the number of predictors. Note that the result varies with different seeds due to randomness of the splitting processes (data, branches). The criterion which uses all predictors for each split has the highest test error being around 15 to 17. It seems that the test error stabilizes when the tree size is approximately 150.

## Exercise 2

Use the bagging approach in order to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important.

```
# Bagging
set.seed(123)
# randomForest() method can be used to bagging if mtry is equal to number of
# all predicting attributes
bag.carseats = randomForest(Sales ~ ., data = trainCAR, mtry = 10, ntree = 500,
                             importance = T)

# Making predictions
bag.pred = predict(bag.carseats, testCAR)

# Test MSE
paste("The test MSE is", mean((testCAR$Sales - bag.pred)^2))

# Importance
paste("Importance of attributes")
importance(bag.carseats)
```

```
[1] "The test MSE is 2.28200717599288"
[1] "Importance of attributes"
      %IncMSE IncNodePurity
CompPrice  26.264150      192.24925
```

Income	9.517596	114.31264
Advertising	13.056754	114.11755
Population	1.195061	81.77791
Price	60.006974	598.03404
ShelveLoc	68.487544	592.00725
Age	23.012641	242.63878
Education	3.101703	66.83355
Urban	2.565957	10.91683
US	2.620459	10.43320

Bagging improves the test MSE to 2.28 and the three most important predictors are ShelveLoc, Price and CompPrice in order of importance.

Use random forests to analyze this data. What test MSE do you obtain? Use the `importance()` function to determine which variables are most important. Describe the effect of `m`, the number of variables considered at each split, on the error rate obtained.


```
# Random forest using squareroot of attributes
bag.carseats = randomForest(Sales~., data=trainCAR, mtry=sqrt(10),
                             ntree=500, importance=T)
bag.pred = predict(bag.carseats, testCAR)
paste("The test MSE is", mean((testCAR$Sales - bag.pred)^2))
paste("Importance of attributes")
importance(bag.carseats)

# Simulating with different m
testMSE = c()
for (m in 1:9){
  set.seed(123)
  bag.carseats = randomForest(Sales~., data=trainCAR, mtry=m, ntree=500)
  bag.pred = predict(bag.carseats, testCAR)
  testMSE[m] = mean((testCAR$Sales - bag.pred)^2)
}
plot(testMSE, type = "b", xlab = "Size of m", ylab = "Test MSE",
     main = "Test MSE for different m")
```

```
[1] "The test MSE is 2.79832577920864"
```

```
[1] "Importance of attributes"
```

	%IncMSE	IncNodePurity
CompPrice	10.9586679	185.05891
Income	5.6921913	171.21686
Advertising	11.7827778	151.55667
Population	1.9851092	140.35235
Price	39.8149820	460.90185
ShelveLoc	47.1279904	446.25433
Age	19.2939832	265.12699
Education	2.7993297	90.39131
Urban	0.2740763	20.84900
US	5.1200974	23.81024



The test MSE increases a bit to 2.71 and the third most important predictor changed into Age.

### Exercises 3 and 4

We now use boosting to predict Salary in the Hitters data set.

(a-b) Remove the observations for whom the salary information is unknown, and then log-transform the salaries. Create a training set consisting of the first 200 observations, and a test set consisting of the remaining observations.

```
# Remove observations with missing values in salary
remove = which(is.na(Hitters$Salary))
HT = Hitters[-remove, ]

# log-transform
HT$Salary = log(HT$Salary)

# Training and test set
train = 1:200
trainHT = HT[train, ]
testHT = HT[-train, ]
```

(c) Perform boosting on the training set with 1,000 trees for a range of values of the shrinkage parameter  $\lambda$ . Produce a plot with different shrinkage values on the x-axis and the corresponding training set MSE on the y-axis.

```
library(gbm)
```

```

set.seed(123)

# Vector of different powers
pows = seq(-10, -0.2, by = 0.1)


# Vector of different lambdas
lambdas = 10^pows
length.lambdas = length(lambdas)

# Empty vectors to store errors
train.errors = rep(NA, length.lambdas)
test.errors = rep(NA, length.lambdas)

# For each value of lambda
for (i in 1:length.lambdas) {
  # Performing boosting for 100 trees with given lambda value
  boost.hitters = gbm(Salary ~ ., data = trainHT, distribution = "gaussian",
    n.trees = 1000, shrinkage = lambdas[i])
  # Storing predicted values
  train.pred = predict(boost.hitters, trainHT, n.trees = 1000)
  test.pred = predict(boost.hitters, testHT, n.trees = 1000)
  # Calculating and storing errors
  train.errors[i] = mean((trainHT$Salary - train.pred)^2)
  test.errors[i] = mean((testHT$Salary - test.pred)^2)
}


# Plotting lambdas against training errors
par(mar = c(5.1, 4.1, 1, 2.1))
plot(lambdas, train.errors, type = "b", xlab = "Shrinkage", ylab = "Train MSE",
  col = 4, pch = 20)

```



(d) Produce a plot with different shrinkage values on the x-axis and the corresponding test set MSE on the y-axis.

```
[1] "The minimum test error for boosting is 0.26 for lambda value of 0.25"
```



As seen above there's a good reason why training error shouldn't be used as a criterion for choosing the final model or model parameters.

(e) Compare the test MSE of boosting to the test MSE that results from applying two of the regression approaches seen in Chapters 3 and 6.

```
# Fitting linear model
lm.fit = lm(Salary ~ ., data = trainHT)

# Forming predictions
lm.pred = predict(lm.fit, testHT)

# Calculating test MSE
paste("The test MSE for linear model is", round(mean((testHT$Salary - lm.pred)^2),
2))

library(glmnet)

set.seed(123)

# Forming model matrices and response vector y
X = model.matrix(Salary ~ ., data = trainHT)
X.test = model.matrix(Salary ~ ., data = testHT)
y = trainHT$Salary

# Fitting lasso
```

```

lasso.cv = cv.glmnet(X, y)
minlambda = lasso.cv$lambda.min

# Storing predictions
lasso.pred = predict(lasso.cv$glmnet.fit, s = minlambda, newx = X.test)

# Test MSE
paste("The test MSE for lasso is", round(mean((testHT$Salary - lasso.pred)^2),
2))

```

```

[1] "The test MSE for linear model is 0.49"
[1] "The test MSE for lasso is 0.47"

```

Both Lasso and OLS model have higher test error rates than boosting.

(f) Which variables appear to be the most important predictors in the boosted model?

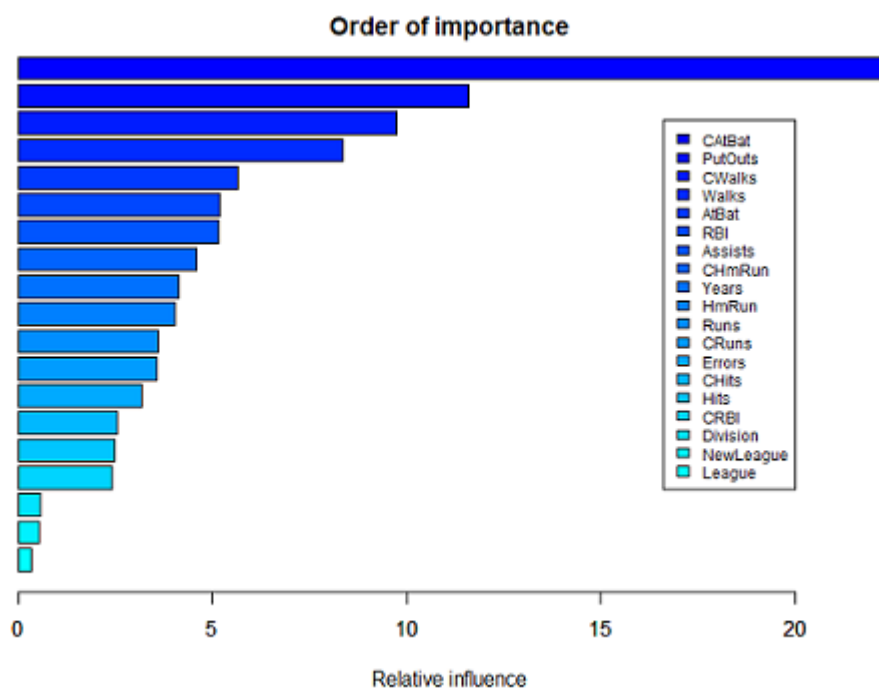
```

# Boosted model which uses the lambda for which the test error minimizes
boost.best = gbm(Salary ~ ., data = trainHT, distribution = "gaussian", n.trees = 1000,
  shrinkage = lambdas[which.min(test.errors)])

head(summary(boost.best, legend.text = T, args.legend = list(cex = 0.8, x = 20,
  y = 20), names.arg = "", main = "Order of importance", xlim = c(0, 20)))

```

	var	rel.inf
CAtBat	CAtBat	15.394543
CWalks	CWalks	11.454057
PutOuts	PutOuts	10.021110
Walks	Walks	8.711869
AtBat	AtBat	6.144511
CRuns	CRuns	5.682345



The three most important predictors in decreasing order of importance are CAtBat, PutOuts and CWalks.

(g) Now apply bagging to the training set. What is the test set MSE for this approach?

```
set.seed(123)

# Bagging mtry = number of all predictors
rf.hitters = randomForest(Salary ~ ., data = trainHT, ntree = 500, mtry = 19)

# Predictions
rf.pred = predict(rf.hitters, testHT)

# Test MSE
paste("The test MSE for bagging is", round(mean((testHT$Salary - rf.pred)^2),
2))
```

```
[1] "The test MSE for bagging is 0.23"
```

The test error for bagging is slightly lower than test MSE for boosting.

## Exercise 5

This question uses the Caravan data set.

(a) Create a training set consisting of the first 1,000 observations, and a test set consisting of the remaining observations.

```
# Checking for missing values
paste("Number of missing values:", sum(is.na(Caravan)))

caravan = Caravan

# Decoding response variable as dummy variable
caravan$Purchase = ifelse(caravan$Purchase == "Yes", 1, 0)

# Training and test set
train = 1:1000
trainCaravan = caravan[train, ]
testCaravan = caravan[-train, ]
```

```
[1] "Number of missing values: 0"
```

(b) Fit a boosting model to the training set with Purchase as the response and the other variables as predictors. Use 1,000 trees, and a shrinkage value of 0.01. Which predictors appear to be the most important?

```
# Boosting (response is binary, so distribution is bernoulli)
set.seed(123)
boostCaravan = gbm(Purchase ~ ., data = trainCaravan, distribution = "bernoulli",
n.trees = 1000, shrinkage = 0.01)
```



```
# Importance
head(summary(boostCaravan, plotit = F))
```

```
      var    rel.inf
PPERSAUT PPERSAUT 14.681380
MKOOPKLA MKOOPKLA  9.244102
MOPLHOOG MOPLHOOG  7.156570
MBERMIDD MBERMIDD  6.311954
PBRAND    PBRAND   5.410268
MGODGE     MGODGE  4.557510
```

The three most important predictors in decreasing order of importance are PPERSAUT, MKOOPKLA and MOPLHOOG

(c) Use the boosting model to predict the response on the test data. Predict that a person will make a purchase if the estimated probability of purchase is greater than 20 %. Form a confusion matrix. What fraction of the people predicted to make a purchase do in fact make one? How does this compare with the results obtained from applying KNN or logistic regression to this data set?

```
# Predictions
boost.prob = predict(boostCaravan, testCaravan, n.trees = 1000, type = "response")

# Storing if probability is > 0.2
boost.pred = ifelse(boost.prob > 0.2, 1, 0)

# Confusion matrix
confm = table(testCaravan$Purchase, boost.pred)
confm

# Actually buy
(confm[2, 2])/(sum(confm[, 2]))
```

```
      boost.pred
      0      1
0 4420  113
1  256   33
[1] 0.2260274
```

Approximately 23 % of customers predicted to make the purchase actually buy the product.

```
library(class)

# Preparing data
caravan = Caravan

# Recommended to standardize data
stand.X = scale(caravan[, -86])
train = 1:1000
train.stan.X = stand.X[train, ]
test.stan.X = stand.X[-train, ]
train.y = caravan[train, 86]
test.y = caravan[-train, 86]
```

```

# kNN
set.seed(123)
kNN = knn(train.stan.X, test.stan.X, train.y, k = 2)

# Confusion matrix
confmt = table(kNN, test.y)
confmt

# Actually buy
(confmt[2, 2])/(sum(confmt[, 2]))

```

```

      test.y
kNN      No  Yes
  No  4246  253
  Yes   287   36
[1] 0.1245675

```

Approximately 12 % of customers predicted to make the purchase actually buy the product which is less than previously got with boosting.